

RAG Assistant - Sample Test Document

This is a throwaway PDF used to smoke-test the ingestion pipeline.

Retrieval-Augmented Generation (RAG) combines a retriever and a generator. Documents are split into chunks, embedded into vectors, and stored in a vector database such as Chroma. At query time, the most relevant chunks are retrieved and passed to the language model as context so it can answer grounded in the source material.

The embeddings here are produced by the nomic-embed-text model served locally through Ollama. Chunks are 512 characters with a 50-character overlap to preserve context across boundaries.

You can replace this file with any real PDF in the docs/ folder.